



# PRAETORIAN

## Symbolic Execution

HOW NOT TO BE ANGR-Y

# \$ id

- ▶ Anthony Weems
- ▶ Staff Engineer at Praetorian
- ▶ Web applications + fun vuln research stuff
- ▶ @amlweems

# Install all the things

- ▶ Follow along: [demo.praetorian.com](https://demo.praetorian.com)
- ▶ Slides available: [demo.praetorian.com/slides](https://demo.praetorian.com/slides)
- ▶ If angr hasn't been installed and you want join the fun, check out <https://github.com/angr/angr> for installation
- ▶ Alternative: Install EpicTreasure for an all-in-one VM (that can also be used in CTFs) <https://github.com/praetorian-inc/epictreasure>



# Quick background


- ▶ Symbolic Execution simply means analyzing a program symbolically to determine what inputs will exercise a given path
- ▶ Layman's terms: How can we traverse from Point A to Point B in a program
- ▶ Accomplished by converting an execution path into a boolean satisfiability problem and throwing the problem at a SAT solver
- ▶ By definition, the solution to these equations is an input that will follow the given execution path

# Quick background


- ▶ Has been a thriving academic research topic since the 1980's
- ▶ The problem was the computational power required for SAT solvers to solve equations in a reasonable amount of time
- ▶ Symbolic execution has been growing in popularity with the advances in publicly accessible SAT solvers (see Microsoft's Z3)
- ▶ Usage of symbolic execution has also been seen in DARPA's recent Cyber Grand Challenge

# Goal: Easy crackme

```
int main(int argc, char** argv) {  
    if(argv[1][0] == 'a') {  
        if(argv[1][1] == 't') {  
            if(argv[1][2] == 'x') {  
                printf("Hooray!");  
            }  
        }  
    }  
}
```



# Goal: Easy crackme




No constraints

&& input[0] == 'a'

&& input[0] != 'a'

# Goal: Easy crackme



No constraints


`&& input[0] == 'a'`

`&& input[0] != 'a'`

`continue`



# Goal: Easy crackme



No constraints


`&& input[0] == 'a'`

`&& input[0] != 'a'`

`continue`

`exit()`

# Goal: Easy crackme




`input[0] == 'a'`

`&& input[1] == 't'`

`&& input[1] != 't'`

# Goal: Easy crackme



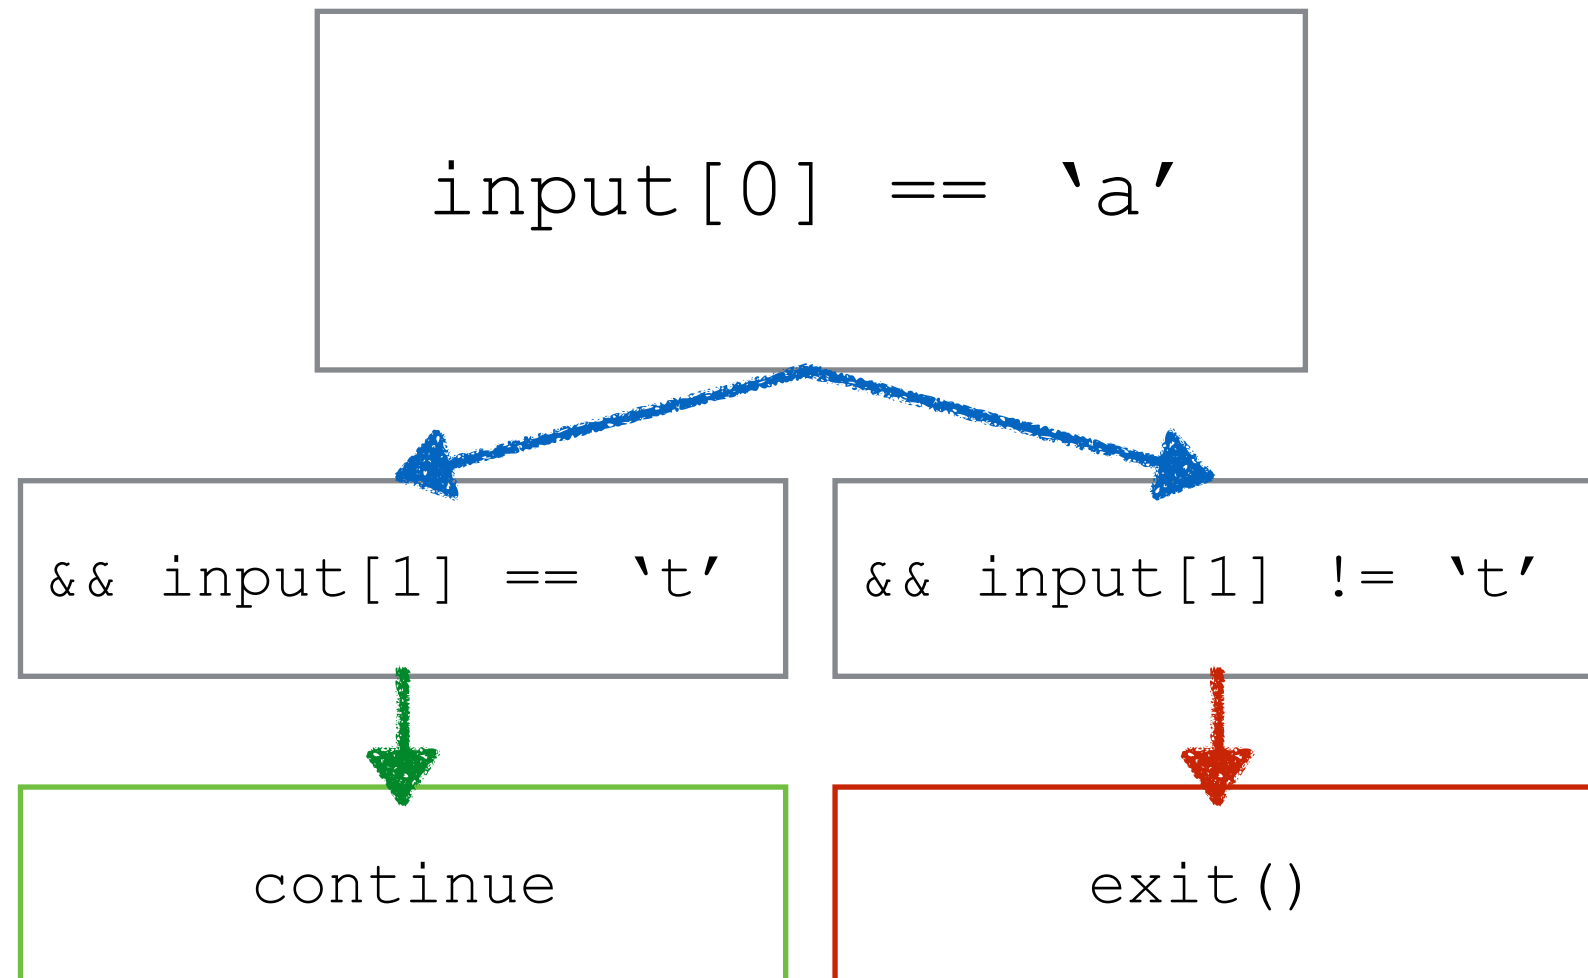

```
input[0] == 'a'
```

```
&& input[1] == 't'
```

```
&& input[1] != 't'
```


```
continue
```

# Goal: Easy crackme





# Goal: Easy crackme




```
input[0] == 'a' &&  
input[1] == 't'
```

```
&& input[2] == 'x'
```

```
&& input[2] != 'x'
```

# Goal: Easy crackme




```
input[0] == 'a' &&  
input[1] == 't'
```

```
&& input[2] == 'x'
```

```
&& input[2] != 'x'
```

```
$$$
```

# Goal: Easy crackme



```
input[0] == 'a' &&  
input[1] == 't'
```


```
&& input[2] == 'x'
```

```
$$$
```

```
&& input[2] != 'x'
```

```
exit()
```

# Goal: Easy crackme



```
input[0] == 'a' &&  
input[1] == 't' &&  
input[2] == 'x'
```

```
print("Hooray!")
```

**Simply solve the found constraints to extract the answer**



# Practical exercise 1

- ▶ We will be using angr as our symbolic engine of choice
- ▶ Developed by researchers from UC Santa Barbara
- ▶ Shellphish CTF team
- ▶ DARPA Cyber Grand Challenge finalists
- ▶ Powered part of their Cyber Reasoning System
- ▶ Python ;-)



SOURCE: [HTTPS://GITHUB.COM/ANGR](https://github.com/angr)

# Practical exercise 1

## ► Quick angr terminology:

- project - binary blob or executable currently being analyzed
- state - an emulated machine state
- path - series of basic blocks representing the current execution flow
- path-group - collection of paths with an easy interface with explore all paths at once

## ► Input sources

- A common problem with symbolic execution is telling the engine what to consider symbolic for analysis
- Common locations for input: stdin, argv, sockets

# Practical exercise 1

- ▶ Basic reverse engineering before starting angr:
  - Included in EpicTreasure is radare2, a command line disassembler
  - Can use IDA Pro, Hopper, or even ObjDump for this portion as well
  - Find the address of the basic block housing the printf to guide angr to the solution we want



[HTTPS://AVATARS3.GITHUBUSERCONTENT.COM/U/917142?V=3&S=400](https://avatars3.githubusercontent.com/u/917142?v=3&s=400)

# Practical exercise 1

- ▶ Quick guide for navigating around radare2:
  - `r2 binary` - Start the binary in radare2
  - `aaa` - Analyze all the things in the binary
  - `afl` - List functions
  - `s main` - Seek to an address for analysis (i.e. main)
  - `VV` - Visual mode
    - ▶ `'hjkl'` or Arrow keys for movement
    - ▶ `q` to quit
  - `?` appended to any commands for lots of usage
    - ▶ `?, a?, af?`



# Practical exercise 1

---

**Locate the basic block address with the printf**

# Practical exercise 1

---

**Wanted address: 0x40057a**

# Practical exercise 1

- ▶ Begin by loading the binary into angr

```
import angr
proj = angr.Project('demo1')
```

- ▶ Create the symbolic input for argv using the provided solver engine

```
import claripy
input = state.se.BVS('input', 3 * 8)
```

- ▶ The BVS is a symbolic bit vector.

Each BVS needs an id ("input") and a length in bits (3 bytes or  $3 * 8$ ).

- ▶ We want to tell angr to set the current state to the entry point with our argv

```
state = proj.factory.entry_state(args=['demo1', input])
```

# Practical exercise 1

- ▶ Create a path group from the entry state

```
pg = proj.factory.path_group(state)
```

- ▶ Tell angr to continue searching paths until we reach the target block

```
pg.explore(find=0x40057a)
```

- ▶ At this point, angr has found a path to reach the given destination block

- ▶ Extract the found state

```
state = pg.found[0].state
```



# Practical exercise 1

- ▶ We can look at the constraints angr found to reach the destination

```
print(state.simplify())
```

```
[<Bool And((input_8_24[23:16] == 97),  
            (input_8_24[15:8] == 116),  
            (input_8_24[7:0] == 120))>]
```

- ▶ Finally, ask angr to solve these constraints for our answer

```
print(state.se.any_str(input))
```

```
atx
```

# Practical exercise 1

```
import angr
import claripy

proj = angr.Project('demo1')
input = claripy.BVS('input', 3 * 8)

state = proj.factory.entry_state(args=[proj.filename, input])
pg = proj.factory.path_group(state)
pg = pg.explore(find=0x40057a)

state = pg.found[0].state
print(state.simplify())

print(state.se.any_str(input))
```

# Exercise 2: Einstein Riddle

- ▶ There are 5 houses in five different colors...
- ▶ The Brit lives in the red house... etc.
- ▶ Lots of boolean logic, prime target for symbolic execution (z3)

	RED	WHITE	YELLOW	BLUE	GREEN	HOUSE	PETS	DRINKS	SMOKES
1	✓	✓	✓	✓	✓	1	✓	✓	✓
2	✓	✓	✓	✓	✓	2	✓	✓	✓
3	✓	✓	✓	✓	✓	3	✓	✓	✓
4	✓	✓	✓	✓	✓	4	✓	✓	✓
5	✓	✓	✓	✓	✓	5	✓	✓	✓
PALL MALL	✓	✓	✓	✓	✓				
DUNHILL	✓	✓	✓	✓	✓				
BLEND	✓	✓	✓	✓	✓				
BLUEMASTER	✓	✓	✓	✓	✓				
PRINCE	✓	✓	✓	✓	✓				
TEA	✓	✓	✓	✓	✓				
COFFEE	✓	✓	✓	✓	✓				
MILK	✓	✓	✓	✓	✓				
BEER	✓	✓	✓	✓	✓				
WATER	✓	✓	✓	✓	✓				
DOG	✓	✓	✓	✓	✓				
BIRD	✓	✓	✓	✓	✓				
CAT	✓	✓	✓	✓	✓				
HORSE	✓	✓	✓	✓	✓				
FISH	✓	✓	✓	✓	✓				
BRITISH	✓	✓	✓	✓	✓				
SWEDISH	✓	✓	✓	✓	✓				
DANISH	✓	✓	✓	✓	✓				
NORWEGIAN	✓	✓	✓	✓	✓				
GERMAN	✓	✓	✓	✓	✓				

green is L of something  
so not 4.5  
white has green to L  
so not 1  
NORWEGIAN and next to blue = 1.5

# Exercise 2: Einstein Riddle

- ▶ `puzzle` program reads binary data from stdin
- ▶ Assigns the data to colors, drinks, nations, etc.
- ▶ Matches against Einstein Riddle rules

```
- $ printf "%025d" | ./puzzle
48:Blue
48:Green
48:Ivory
48:Red
48:Yellow
...
womp womp... try again
```

# Exercise 2: Einstein Riddle Input

- ▶ Byte values 1-5 represents house number of item
  - e.g. Milk in middle house represented by 3
  - e.g. Norwegian in first house represented by 1
- ▶ Specify items in order (found inside binary, but not important)
  - ```
$ echo -ne "\x01\x02\x03\x04\x05" | ./puzzle
1:Blue
2:Green
3:Ivory
4:Red
5:Yellow
...
womp womp... try again
```



# Exercise 2: Einstein Riddle Constraints

Libraries can be nasty, we can disable them and use a simulated libc

```
opts = {'auto_load_libs': False}
proj = angr.Project('puzzle', load_options=opts)
st    = proj.factory.entry_state(args=['puzzle'])

for _ in range(25):
    # works like reading from a fd, moves seek head
    e = st.posix.files[0].read_from(1)
```

**Add constraints to SAT problem**

```
st.add_constraints(e >= 1)
st.add_constraints(e <= 5)
```

```
e = st.posix.files[0].read_from(1)
st.add_constraints(e == 0)
```

**Set the length of stdin**

```
st.posix.files[0].seek(0)
st.posix.files[0].length = 25
```

# Exercise 2: Einstein Riddle

**Make puzzle print “congratulations!”**

**Who drinks water? Who owns the zebra?**

# Exercise 3

---

**Can we only solve reverse engineering problems?**

# Exercise 3: Stack overflow

```
#include <string.h>
#include <stdio.h>

void overflow_me(){
    char name[24];
    printf("Welcome.. what is your name?\n");
    read(0, name, 80);
    return;
}

int main(int argc, char** argv){
    char vuln[32];
    printf("Password protected. Enter password:\n");

    read(0, vuln, 32);
    if(strstr(vuln, "badpassword") == vuln)
        overflow_me();
    else
        printf("Wrong password\n");

    return 0;
}
```

# Stack overflow review

```
void overflow_me() {  
    char name[24];  
    read(0, name, 80);  
    return;  
}
```

Begin with the stack frame from the vulnerable function.






# Stack overflow review

```
void overflow_me() {  
    char name[24];  
    read(0, name, 80);  
    return;  
}
```


We read beyond the bounds of the char array, overwriting the saved EIP.



# Stack overflow review

```
void overflow_me() {  
    char name[24];  
    read(0, name, 80);  
    return;  
}
```

On return, we **SEGFault** because our EIP is now at an invalid address.



# Exercise 3: Stack overflow

- ▶ By default, angr discards any unconstrained paths. In this example, we are exactly looking for unconstrained paths:
  - `pg = proj.factory.path_group(state, save_unconstrained=True)`
- ▶ We can explore all paths until we have a constrained path. Instead of `explore`, we will manually `step` each path until the number of unconstrained paths is greater than 0.
  - `pg.step(until=lambda x: len(x.unconstrained) > 1)`
- ▶ After each path is stepped forward once, it is checked for unconstrained paths. If we found one, `step` will return and hand us our wanted path.

# Exercise 3: Stack overflow

- ▶ Grab our path from the unconstrained array.
  - `state = pg.unconstrained[0].state`
- ▶ We now can simply set a constraint such that the state's Instruction Pointer equals whatever we want.
  - `crash_ip = claripy.BVV(int('deadbeefcafebabe', 16), 8 * 8)`
  - `state.se.add(state.regs.ip == crash_ip)`
- ▶ Because we added the IP constraint, any solved solution will by definition crash with `IP == deadbeefcafebabe`
  - `payload = state.posix.dumps(0)`

# Exercise 3: Stack overflow

---

## Demo3



# Exercise 3: Stack overflow

```
$ xxd solution
00000000: 6261 6470 6173 7377 6f72 6400 0000 0000  badpassword.....
00000010: 0000 0000 0000 0000 0000 0000 0000 0000  .....
00000020: 0000 0000 0000 0000 0000 0000 0000 0000  .....
00000030: 0000 0000 0000 0000 0000 0000 0000 0000  .....
00000040: 0000 0000 0000 0000 beba fecb efbe adde  .....
00000050: 0000 0000 0000 0000 0000 0000 0000 0000  .....
00000060: 0000 0000 0000 0000 0000 0000 0000 0000  .....
```



# Further Reading

## ▶ Papers

- [http://www.internetsociety.org/sites/default/files/11\\_1\\_2.pdf](http://www.internetsociety.org/sites/default/files/11_1_2.pdf)

## ▶ angr docs

- <https://github.com/angr/angr-doc>
- Loads of examples and detailed explanation of API

## ▶ Cyber Grand Challenge

- <http://www.cybergrandchallenge.com/>
- <https://github.com/CyberGrandChallenge/samples>
  - ▶ cqe-challenges contains lots of juicy exploitable binaries



# PRAETORIAN

## Symbolic Execution Workshop

HOW NOT TO BE ANGR-Y